

# DB - HW 11

## 15.2

Consider the bank database of Figure 15.14, where the primary keys are underlined, and the following SQL query:

```
SELECT  T.branch_name
FROM    branch T, branch S
WHERE   T.assets > S.assets AND S.branch_city = "Brooklyn"
```

Write an efficient relational-algebra expression that is equivalent to this query.

Justify your answer.

Answer:

$$\Pi_{T.branch\_name}((\Pi_{T.branch\_name, assets}(\rho_T(branch))) \bowtie_{T.assets > S.assets} (\Pi_{assets}(\sigma_{branch\_city='Brooklyn'}(\rho_S(branch)))))$$

We do the selection  $\sigma_{branch\_city='Brooklyn'}$  before doing the Theta Join, which eliminated irrelevant tuples and thus significantly reduced the number of data involved in this process.

## 15.3

Let relations  $r_1(A,B,C)$  and  $r_2(C,D,E)$  have the following properties:  $r_1$  has 20,000 tuples,  $r_2$  has 45,000 tuples, 25 tuples of  $r_1$  fit on one block, and 30 tuples of  $r_2$  fit on one block. Estimate the number of block transfer and seeks required using each of the following join strategies for  $r_1 \bowtie r_2$ :

a. Nested-loop join.

b. Block nested-loop join.

c. Merge join.

d. Hash join.

Answer:

$r_1$  needs  $20000 \div 25 = 800$  blocks

$r_2$  needs  $45000 \div 30 = 1500$  blocks

Assuming the memory has N pages:

a. Nested-loop Join

- Use  $r_1$  as outer relation:  $20000 * 1500 + 800 = 30000800$
- Use  $r_2$  as outer relation:  $45000 * 800 + 1500 = 36001500$

b. Block nested-loop join

- Use  $r_1$  as outer relation:  $\lceil \frac{800}{N-1} \rceil * 1500 + 800$
- Use  $r_2$  as outer relation:  $\lceil \frac{1500}{N-1} \rceil * 800 + 1500$

c. Merge join

If both of  $r_1$  and  $r_2$  are not sorted on the join key, the cost of sorting would be:

$$1500 * (2 * \lceil \log_{N-1}(\frac{1500}{N}) \rceil + 2) + 800 * (2 * \lceil \log_{N-1}(\frac{800}{N}) \rceil + 2)$$

Assuming all the tuples possessing the same value for join attributes fit in memory, the total cost would be:

$$1500 * (2 * \lceil \log_{N-1}(\frac{1500}{N}) \rceil + 2) + 800 * (2 * \lceil \log_{N-1}(\frac{800}{N}) \rceil + 2) + 1500 + 800$$

d. Hash join

Assuming there is no flow, using  $r_1$  as build relation and  $r_2$  as probe relation:

If  $N > \frac{800}{N}$ :

As there is no need to do recursive partitioning, the cost would be  $3 * 1500 + 800 = 6900$

Else:

The cost would be  $2 * (1500 + 800) * \lceil \log_{N-1}(800) - 1 \rceil + 1500 + 800$

## 15.6

Consider the bank database of Figure 15.14, where the primary keys are underlined. Suppose that a  $B^+$ -tree index on branch\_city is available on relation *branch*, and that no other index is available. List different ways to handle the following selections that involve negation:

- $\sigma_{\neg(\text{branch\_city} < \text{'Brooklyn'})}(\text{branch})$
- $\sigma_{\neg(\text{branch\_city} = \text{'Brooklyn'})}(\text{branch})$
- $\sigma_{\neg(\text{branch\_city} < \text{'Brooklyn'} \vee \text{assets} < 5000)}(\text{branch})$

Answer:

$$a. \because \neg(\text{branch\_city} < \text{'Brooklyn'}) \iff \text{branch\_city} \geq \text{'Brooklyn'}$$

Thus, we just need to locate the first tuple meeting the requirement:  $\text{record.branch\_city} = \text{'Brooklyn'}$  by using the index, and then retrieve all the tuple from this position to the end by following the pointer chains.

$$b. \because \neg(\text{branch\_city} = \text{'Brooklyn'}) \iff \text{branch\_city} \neq \text{'Brooklyn'}$$

In this case, scan the whole file sequentially performs better than using the index. Thus, we can simply scan the whole file in sequential order and retrieve all the tuple whose value of field *branch\_city* is anything but 'Brooklyn'.

c.

$$\because \neg(\text{branch\_city} < \text{'Brooklyn'} \vee \text{assets} < 5000) \iff \text{branch\_city} \geq \text{'Brooklyn'} \wedge \text{assets} \geq 5000$$

Thus, like what we did in a), we can find the first tuple whose value of *branch\_city* equals to 'Brooklyn' by using the index of *branch\_city*, and follow the pointer chain to retrieve every tuple whose value of attribute *branch\_city* is bigger than or equal to 'Brooklyn'. Additionally, we should check whether the value of attribute *assets* we visit is no less than 5000.

Or we can find the first tuple whose value of attribute *assets* is no less than 5000 and retrieve all the tuples till the end of file. Besides, checking whether the value of attribute *branch\_city* is bigger than or equal to 'Brooklyn'.

## 15.20

**Estimate the number of block transfers and seeks required by your solution to Exercise 15.19 for  $r_1 \bowtie r_2$  are as defined in Exercise 15.3**

Answer:

In Exercise 15.3 we have already calculated that  $r_1$  and  $r_2$  takes 800 and 1500 blocks respectively.

Assuming there are  $n$  pointers in each leaf block of an index, then:

index of  $r_1$  needs  $\lceil \frac{20000}{n} \rceil$  leaf blocks, and index of  $r_2$  needs  $\lceil \frac{45000}{n} \rceil$  leaf blocks.

Thus, without output, the merge sort would take  $\lceil \frac{20000}{n} \rceil + \lceil \frac{45000}{n} \rceil$  times of access.

The number of output tuples =  $\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil$ , each of them would take 2 pointers.

The number of blocks for output of join =  $\lceil \frac{\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil}{n/2} \rceil = \lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil$

The total disc block accesses of the join =  $\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil + \lceil \frac{20000}{n} \rceil + \lceil \frac{45000}{n} \rceil$

Assuming the memory has  $M$  pages, when replacing the pointers by actual tuples, we needs:

$\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil * (2 + 2 * \log_{M-1}(\frac{\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil}{M}))$  disc block accesses for the first sorting.

Besides, there are  $\min(800, \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil)$  blocks in  $r_1$  whose pointer should be replaced.

Assuming that in 1 disk block, there are  $K_1$  pairs of (tuple\_of\_ $r_1$ , ptr\_to\_ $r_2$ ), thus the intermediate result would take  $\lceil \frac{\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil}{K_1} \rceil$  blocks.

The first pass of replacement will cost:

$$\begin{aligned} & \lceil \frac{2}{n} * \lceil \frac{20000 * 45000}{\max(V(C, r_1), V(C, r_2))} \rceil \rceil * (2 + 2 * \log_{M-1}(\frac{\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C, r_1), V(C, r_2))} \rceil \rceil}{M})) + \lceil \frac{2}{n} * \lceil \frac{20000 * 45000}{\max(V(C, r_1), V(C, r_2))} \rceil \rceil \\ & + \min(800, \lceil \frac{20000 * 45000}{\max(V(C, r_1), V(C, r_2))} \rceil) + \lceil \frac{\lceil \frac{20000*45000}{\max(V(C, r_1), V(C, r_2))} \rceil}{K_1} \rceil \quad \text{times of block accesses} \end{aligned}$$

Likewise, assuming that there are  $K_2$  pairs of (tuple\_of\_ $r_2$ , ptr\_to\_ $r_1$ ) in 1 block,.

The second pass of replacement will cost:

$$\begin{aligned} & \lceil \frac{\lceil \frac{20000*45000}{\max(V(C, r_1), V(C, r_2))} \rceil}{K_1} \rceil * (2 + 2 * \log_{M-1}(\frac{\lceil \frac{\lceil \frac{20000*45000}{\max(V(C, r_1), V(C, r_2))} \rceil}{K_1} \rceil}{M})) \\ & + \lceil \frac{\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C, r_1), V(C, r_2))} \rceil \rceil}{K_1} \rceil + \min(1500, \lceil \frac{20000 * 45000}{\max(V(C, r_1), V(C, r_2))} \rceil) \end{aligned}$$

By adding up all the three parts, we can have the total number of disc access for the join:

$$\begin{aligned}
& \lceil \frac{20000}{n} \rceil + \lceil \frac{45000}{n} \rceil + 2 * \lceil \frac{\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil}{K_1} \rceil * (2 + \log_{M-1}(\frac{\lceil \frac{\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil}{K_1} \rceil}{M})) \\
& + 2 * \lceil \frac{2}{n} * \lceil \frac{20000 * 45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil * (2 + \log_{M-1}(\frac{\lceil \frac{2}{n} * \lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil \rceil}{M})) \\
& + \min(1500, \lceil \frac{20000 * 45000}{\max(V(C,r_1), V(C,r_2))} \rceil) + \min(800, \lceil \frac{20000 * 45000}{\max(V(C,r_1), V(C,r_2))} \rceil)
\end{aligned}$$

And the number of output pages =  $\lceil \frac{\lceil \frac{20000*45000}{\max(V(C,r_1), V(C,r_2))} \rceil}{K_2} \rceil$